

CSE 232: Programming Assignment 2

Socket Programming with Performance Analysis

Himanshu Kumar(2022215), Himanshu Raj2022216

Q.1. TCP multithreaded client-server connection

1) The server sets up a TCP socket, binds a buffer space to receive requests and listens for client connections.

```
// Create socket file descriptor
if ((server_fd = socket(AF_INET, SOCK_STREAM, 0)) == 0) {
    perror("Socket failed");
    exit(EXIT_FAILURE);
}

address.sin_family = AF_INET;
address.sin_addr.s_addr = INADDR_ANY;
address.sin_port = htons(PORT);

// Bind the socket to the address
if (bind(server_fd, (struct sockaddr*)&address, sizeof(address)) < 0) {
    perror("Bind failed");
    exit(EXIT_FAILURE);
}

// Listen for incoming connections
if (listen(server_fd, 3) < 0) {
    perror("Listen failed");
    exit(EXIT_FAILURE);
}
```

2) The server accepts the client connection and creates a new thread that continues to process the client connection. The original server socket continues to listen on the same listening port for newer incoming client connections.

```
// Continuously accept incoming connections and handle them in separate threads
while (1) {
    new_socket = accept(server_fd, (struct sockaddr*)&address, (socklen_t*)&addrlen);
    if (new_socket < 0) {
        perror("Accept failed");
        exit(EXIT_FAILURE);
    }

    // Allocate memory for socket descriptor and pass it to the thread
    int* socket_desc = malloc(sizeof(int));
    *socket_desc = new_socket;

    // Create a new thread to handle the client connection
    pthread_t thread_id;
    if (pthread_create(&thread_id, NULL, handle_client, (void*)socket_desc) != 0) {
        perror("Thread creation failed");
        free(socket_desc);
        exit(EXIT_FAILURE);
    }
}
```

3) The client creates n concurrent client connection requests to the server, where n is a program argument.

```
int num_clients = atoi(argv[1]);
pthread_t threads[num_clients];

// Create multiple threads for concurrent client connections
for (int i = 0; i < num_clients; i++) {
    pthread_create(&threads[i], NULL, client_task, NULL);
}

// Wait for all threads to complete
for (int i = 0; i < num_clients; i++) {
    pthread_join(threads[i], NULL);
}
```

4) After the client connection is established, the client sends a request to the server to get information about the server's top TWO CPU-consuming processes.

```
const char* request = "Request: Top CPU Processes";

if ((sock = socket(AF_INET, SOCK_STREAM, 0)) < 0) {
    perror("Socket creation error");
    return NULL;
}

serv_addr.sin_family = AF_INET;
serv_addr.sin_port = htons(PORT);

if (inet_pton(AF_INET, "127.0.0.1", &serv_addr.sin_addr) <= 0) {
    perror("Invalid address/ Address not supported");
    return NULL;
}

if (connect(sock, (struct sockaddr*)&serv_addr, sizeof(serv_addr)) < 0) {
    perror("Connection Failed");
    return NULL;
}
```

5) The server sends the information about the top CPU-consuming process (using proc system call) to the client.

```
// Open the /proc directory
if ((dir = opendir("/proc")) == NULL) {
    perror("Failed to open /proc");
    return;
}

while ((entry = readdir(dir)) != NULL) {
    if (entry->d_type == DT_DIR && atoi(entry->d_name) > 0) {
        // Build path to /proc/[pid]/stat
        snprintf(path, sizeof(path), "/proc/%s/stat", entry->d_name);

        // Open the /proc/[pid]/stat file
        stat_file = fopen(path, "r");
        if (stat_file) {
            fscanf(stat_file, "%d %s %c %d %d %d %d %d %d %u %u %u %u %u %u %d %d",
                &pid, process_name, &user_time, &kernel_time);
            total_time = user_time + kernel_time;
            fclose(stat_file);
        }
    }
}
```

6) The client prints this information & closes the connection.

```
// Read the server's response (information about top CPU-consuming processes)
read(sock, buffer, 1024);
printf("Client: Server response:\n%s\n", buffer);

close(sock);
return NULL;
```

Q.2. Performance Analysis

a) Single-threaded TCP client-server

In the single-threaded model, the server can only handle one client connection at a time. All other client requests are queued until the current request is completed.

```

• iiitd@iiitd-ThinkCentre-M70s-Gen-3:~/Desktop/cn$ sudo taskset -c 1 perf stat ./single_thread_client
Client: Request sent to server
Client: Server response:
Top 2 CPU-consuming processes:
1. Process: (systemd), PID: 2221, CPU time: 15228 ticks
2. Process: (gnome-shell), PID: 2470, CPU time: 8764 ticks

Performance counter stats for './single_thread_client':

    0.37 msec task-clock                #    0.156 CPUs utilized
         2    context-switches          #    5.397 K/sec
         0    cpu-migrations            #    0.000 /sec
        60    page-faults               #   161.920 K/sec
<not counted>    cpu_atom/cycles/              (0.00%)
    1,646,425    cpu_core/cycles/              #    4.443 GHz
<not counted>    cpu_atom/instructions/          (0.00%)
    1,281,034    cpu_core/instructions/          (0.00%)
<not counted>    cpu_atom/branches/              (0.00%)
    228,609     cpu_core/branches/              #   616.938 M/sec
<not counted>    cpu_atom/branch-misses/          (0.00%)
         8,596    cpu_core/branch-misses/
    TopdownL1 (cpu_core)                #    21.1 % tma_backend_bound
                                         #     9.9 % tma_bad_speculation
                                         #    51.8 % tma_frontend_bound
                                         #    17.2 % tma_retiring

    0.002372328 seconds time elapsed

    0.000000000 seconds user
    0.000570000 seconds sys

```

Single-threaded client connecting to single-threaded server

```

• iiitd@iiitd-ThinkCentre-M70s-Gen-3:~/Desktop/cn$ sudo taskset -c 0 perf stat ./single_thread_server
Server is listening on port 8005
New connection: socket fd is 4, ip is : 127.0.0.1, port: 47634
Adding to list of sockets as 0
Message received from client: Requesting top CPU processes
Response sent to client
Host disconnected, ip 127.0.0.1, port 47634
^C./single_thread_server: Interrupt

Performance counter stats for './single_thread_server':

    2.73 msec task-clock                #    0.001 CPUs utilized
         3    context-switches          #    1.100 K/sec
         0    cpu-migrations            #    0.000 /sec
        65    page-faults               #   23.842 K/sec
<not counted>    cpu_atom/cycles/              (0.00%)
    10,695,058    cpu_core/cycles/              #    3.923 GHz
<not counted>    cpu_atom/instructions/          (0.00%)
    17,303,288    cpu_core/instructions/          (0.00%)
<not counted>    cpu_atom/branches/              (0.00%)
    3,182,441     cpu_core/branches/              #    1.167 G/sec
<not counted>    cpu_atom/branch-misses/          (0.00%)
        21,013    cpu_core/branch-misses/
    TopdownL1 (cpu_core)                #    39.6 % tma_backend_bound
                                         #     4.0 % tma_bad_speculation
                                         #    28.2 % tma_frontend_bound
                                         #    28.2 % tma_retiring

    2.866608647 seconds time elapsed

    0.001062000 seconds user
    0.002125000 seconds sys

```

Single-threaded client connecting to **single-threaded server**

```

Performance counter stats for './multi_thread_client 10':

    1.35 msec task-clock                #    0.001 CPUs utilized
      21      context-switches          #   15.540 K/sec
       0      cpu-migrations            #    0.000 /sec
      85      page-faults               #   62.900 K/sec
<not counted>    cpu_atom/cycles/              (0.00%)
  4,762,339    cpu_core/cycles/                #    3.524 GHz
<not counted>    cpu_atom/instructions/          (0.00%)
  4,028,875    cpu_core/instructions/            (0.00%)
<not counted>    cpu_atom/branches/              (0.00%)
   729,593    cpu_core/branches/                #   539.898 M/sec
<not counted>    cpu_atom/branch-misses/          (0.00%)
   17,448    cpu_core/branch-misses/
  TopdownL1 (cpu_core)                #   25.2 % tma_backend_bound
                                           #    7.1 % tma_bad_speculation
                                           #   47.3 % tma_frontend_bound
                                           #   20.3 % tma_retiring

  1.047452774 seconds time elapsed

  0.000000000 seconds user
  0.001553000 seconds sys

```

Multi-threaded (n=10) client connecting to single-threaded server

```

Performance counter stats for './single_thread_server':

   20.52 msec task-clock                #    0.002 CPUs utilized
      3      context-switches          #   146.218 /sec
       0      cpu-migrations            #    0.000 /sec
      64      page-faults               #    3.119 K/sec
<not counted>    cpu_atom/cycles/              (0.00%)
  72,927,439    cpu_core/cycles/                #    3.554 GHz
<not counted>    cpu_atom/instructions/          (0.00%)
  161,356,000    cpu_core/instructions/            (0.00%)
<not counted>    cpu_atom/branches/              (0.00%)
   29,702,089    cpu_core/branches/                #    1.448 G/sec
<not counted>    cpu_atom/branch-misses/          (0.00%)
   115,970    cpu_core/branch-misses/
  TopdownL1 (cpu_core)                #   24.3 % tma_backend_bound
                                           #    3.6 % tma_bad_speculation
                                           #   33.7 % tma_frontend_bound
                                           #   38.4 % tma_retiring

  8.285303467 seconds time elapsed

  0.004209000 seconds user
  0.016839000 seconds sys

```

Multi-threaded (n=10) client connecting to **single-threaded** server

CPU clocks: Relatively low due to sequential processing. The CPU is not highly loaded.

Cache misses: Low, as the data size is small, and there's no multi-threading.

Context switches: Minimal context switches since only one thread is active.

Latency: Each connection waits until the server finishes the previous request, leading to increased latency for later clients.

```

Performance counter stats for './multi_thread_client 25':

      3.62 msec task-clock                #    0.001 CPUs utilized
         57      context-switches        #   15.738 K/sec
          0      cpu-migrations           #    0.000 /sec
        115      page-faults             #   31.751 K/sec
<not counted>    cpu_atom/cycles/                (0.00%)
    9,648,396    cpu_core/cycles/                #    2.664 GHz
<not counted>    cpu_atom/instructions/           (0.00%)
    8,110,344    cpu_core/instructions/           (0.00%)
<not counted>    cpu_atom/branches/               (0.00%)
    1,472,617    cpu_core/branches/              #   406.585 M/sec
<not counted>    cpu_atom/branch-misses/          (0.00%)
    28,336       cpu_core/branch-misses/
TopdownL1 (cpu_core)  #    26.0 % tma_backend_bound
                      #    5.6 % tma_bad_speculation
                      #   46.9 % tma_frontend_bound
                      #   21.4 % tma_retiring

2.722816221 seconds time elapsed

0.000000000 seconds user
0.003681000 seconds sys

```

Multi-threaded (n=25) client connecting to single-threaded server

```

Performance counter stats for './single_thread_server':

      80.61 msec task-clock                #    0.012 CPUs utilized
         11      context-switches        #   136.457 /sec
          0      cpu-migrations           #    0.000 /sec
         66      page-faults             #   818.742 /sec
<not counted>    cpu_atom/cycles/                (0.00%)
   189,224,239    cpu_core/cycles/                #    2.347 GHz
<not counted>    cpu_atom/instructions/           (0.00%)
   410,769,061    cpu_core/instructions/           (0.00%)
<not counted>    cpu_atom/branches/               (0.00%)
    75,622,364    cpu_core/branches/              #   938.109 M/sec
<not counted>    cpu_atom/branch-misses/          (0.00%)
    291,227       cpu_core/branch-misses/
TopdownL1 (cpu_core)  #    24.6 % tma_backend_bound
                      #    3.6 % tma_bad_speculation
                      #   33.5 % tma_frontend_bound
                      #   38.3 % tma_retiring

6.718000548 seconds time elapsed

0.017513000 seconds user
0.063871000 seconds sys

```

Multi-threaded (n=25) client connecting to **single-threaded** server

CPU clocks: Increased due to a higher number of requests, but still sequential.

Cache misses: Slightly increased but still low due to single-threaded operation.

Context switches: Very low (almost negligible).

Latency: The time to serve all clients grows linearly with the number of clients since the server handles each request one at a time.

CPU utilization: Underutilized, since the server could handle more work but is restricted to sequential processing.

b) Concurrent TCP client-server

In the concurrent model, the server uses multiple threads (one per client), so multiple client requests are handled simultaneously.

```
Performance counter stats for './multi_thread_client 10':

    1.24 msec task-clock                #    0.075 CPUs utilized
      20      context-switches          #   16.129 K/sec
       0      cpu-migrations            #    0.000 /sec
      84      page-faults               #   67.742 K/sec
<not counted>  cpu_atom/cycles/          (0.00%)
  5,086,561    cpu_core/cycles/          #    4.102 GHz
<not counted>  cpu_atom/instructions/    (0.00%)
  4,473,604    cpu_core/instructions/
<not counted>  cpu_atom/branches/        (0.00%)
  802,591      cpu_core/branches/        #   647.249 M/sec
<not counted>  cpu_atom/branch-misses/   (0.00%)
  16,844      cpu_core/branch-misses/
  TopdownL1 (cpu_core)                #   24.6 % tma_backend_bound
                                           #    7.1 % tma_bad_speculation
                                           #   46.0 % tma_frontend_bound
                                           #   22.4 % tma_retiring

 0.016521403 seconds time elapsed

 0.000000000 seconds user
 0.001477000 seconds sys
```

Multi-threaded (n=10) client connecting to multi-threaded server

```
Performance counter stats for './multi_thread_server':

   17.35 msec task-clock                #    0.002 CPUs utilized
      2      context-switches          #   115.291 /sec
       0      cpu-migrations            #    0.000 /sec
     102      page-faults               #    5.880 K/sec
<not counted>  cpu_atom/cycles/          (0.00%)
  74,061,248    cpu_core/cycles/          #    4.269 GHz
<not counted>  cpu_atom/instructions/    (0.00%)
 165,289,392    cpu_core/instructions/
<not counted>  cpu_atom/branches/        (0.00%)
 30,485,779     cpu_core/branches/        #    1.757 G/sec
<not counted>  cpu_atom/branch-misses/   (0.00%)
  108,352      cpu_core/branch-misses/
  TopdownL1 (cpu_core)                #   24.1 % tma_backend_bound
                                           #    3.4 % tma_bad_speculation
                                           #   33.7 % tma_frontend_bound
                                           #   38.8 % tma_retiring

10.080972287 seconds time elapsed

 0.001955000 seconds user
 0.015647000 seconds sys
```

Multi-threaded (n=10) client connecting to multi-threaded server

CPU clocks: Increased due to the need to manage multiple threads.

Cache misses: Higher compared to single-threaded due to the need to context switch between threads.

Context switches: Noticeably higher since each client has its own thread, resulting in multiple context switches as the OS manages thread scheduling.

Latency: Significantly lower per client, as clients are served in parallel.

```

Performance counter stats for './multi_thread_client 25':

    2.59 msec task-clock                #    0.059 CPUs utilized
      36      context-switches          #   13.887 K/sec
       0      cpu-migrations             #    0.000 /sec
     115      page-faults               #   44.360 K/sec
<not counted>      cpu_atom/cycles/          (0.00%)
10,483,624      cpu_core/cycles/          #    4.044 GHz
<not counted>      cpu_atom/instructions/    (0.00%)
  9,106,087      cpu_core/instructions/
<not counted>      cpu_atom/branches/        (0.00%)
 1,634,162      cpu_core/branches/          #   630.364 M/sec
<not counted>      cpu_atom/branch-misses/   (0.00%)
    23,130      cpu_core/branch-misses/
TopdownL1 (cpu_core)                #   26.7 % tma_backend_bound
                                           #    5.2 % tma_bad_speculation
                                           #   43.1 % tma_frontend_bound
                                           #   24.9 % tma_retiring

0.044012776 seconds time elapsed

0.000000000 seconds user
0.002302000 seconds sys

```

Multi-threaded (n=25) client connecting to multi-threaded server

```

Performance counter stats for './multi_thread_server':

    44.19 msec task-clock                #    0.005 CPUs utilized
       8      context-switches          #   181.041 /sec
       0      cpu-migrations             #    0.000 /sec
     164      page-faults               #    3.711 K/sec
<not counted>      cpu_atom/cycles/          (0.00%)
195,075,545      cpu_core/cycles/          #    4.415 GHz
<not counted>      cpu_atom/instructions/    (0.00%)
409,974,123      cpu_core/instructions/
<not counted>      cpu_atom/branches/        (0.00%)
75,615,642      cpu_core/branches/          #    1.711 G/sec
<not counted>      cpu_atom/branch-misses/   (0.00%)
    249,062      cpu_core/branch-misses/
TopdownL1 (cpu_core)                #   28.0 % tma_backend_bound
                                           #    3.1 % tma_bad_speculation
                                           #   32.4 % tma_frontend_bound
                                           #   36.5 % tma_retiring

8.268447248 seconds time elapsed

0.008893000 seconds user
0.034586000 seconds sys

```

Multi-threaded (n=25) client connecting to multi-threaded server


```

Performance counter stats for './multi_thread_client 50':

    5.16 msec task-clock                #    0.005 CPUs utilized
      70      context-switches          #   13.556 K/sec
      0      cpu-migrations             #    0.000 /sec
     167      page-faults              #   32.340 K/sec
<not counted>      cpu_atom/cycles/                (0.00%)
 19,281,042      cpu_core/cycles/                #    3.734 GHz
<not counted>      cpu_atom/instructions/            (0.00%)
 16,988,947      cpu_core/instructions/            (0.00%)
<not counted>      cpu_atom/branches/                (0.00%)
  3,048,042      cpu_core/branches/                #   590.269 M/sec
<not counted>      cpu_atom/branch-misses/            (0.00%)
    36,186      cpu_core/branch-misses/
    TopdownL1 (cpu_core)                #   28.2 % tma_backend_bound
                                           #    4.7 % tma_bad_speculation
                                           #   41.0 % tma_frontend_bound
                                           #   26.1 % tma_retiring

    1.056901312 seconds time elapsed

    0.000000000 seconds user
    0.004187000 seconds sys

```

Multi-threaded (n=50) client connecting to multi-threaded server

```

Performance counter stats for './multi_thread_server':

    96.72 msec task-clock                #    0.010 CPUs utilized
      16      context-switches          #   165.422 /sec
      0      cpu-migrations             #    0.000 /sec
     248      page-faults              #    2.564 K/sec
<not counted>      cpu_atom/cycles/                (0.00%)
 368,876,278      cpu_core/cycles/                #    3.814 GHz
<not counted>      cpu_atom/instructions/            (0.00%)
 820,371,473      cpu_core/instructions/            (0.00%)
<not counted>      cpu_atom/branches/                (0.00%)
 151,303,769      cpu_core/branches/                #    1.564 G/sec
<not counted>      cpu_atom/branch-misses/            (0.00%)
    476,226      cpu_core/branch-misses/
    TopdownL1 (cpu_core)                #   24.7 % tma_backend_bound
                                           #    3.1 % tma_bad_speculation
                                           #   33.7 % tma_frontend_bound
                                           #   38.6 % tma_retiring

    10.030677364 seconds time elapsed

    0.015420000 seconds user
    0.078938000 seconds sys

```

Multi-threaded (n=50) client connecting to **multi-threaded server**

CPU clocks: Higher as more threads are created.

Cache misses: Increases with the number of clients due to thread contention for resources.

Context switches: Very high as more threads lead to frequent context switching.

Latency: Remains low for each client as long as the server is not overwhelmed.

CPU utilization: High CPU usage due to thread management, but more efficient utilization of available CPU resources.

c) TCP client-server using “select”

```
Performance counter stats for './multi_thread_client 10':

    1.59 msec task-clock                #    0.002 CPUs utilized
      21      context-switches          #   13.186 K/sec
       0      cpu-migrations            #    0.000 /sec
      83      page-faults              #   52.116 K/sec
<not counted>  cpu_atom/cycles/                (0.00%)
  4,719,230    cpu_core/cycles/                #    2.963 GHz
<not counted>  cpu_atom/instructions/           (0.00%)
  4,031,820    cpu_core/instructions/           (0.00%)
<not counted>  cpu_atom/branches/              (0.00%)
   730,373    cpu_core/branches/              #   458.608 M/sec
<not counted>  cpu_atom/branch-misses/         (0.00%)
    17,737    cpu_core/branch-misses/
  TopdownL1 (cpu_core)                #   22.0 % tma_backend_bound
                                           #    7.6 % tma_bad_speculation
                                           #   48.9 % tma_frontend_bound
                                           #   21.5 % tma_retiring

  1.054229798 seconds time elapsed

  0.001437000 seconds user
  0.000233000 seconds sys
```

Multi-threaded (n=10) client connecting to server using select

```
Performance counter stats for './select_server':

    21.01 msec task-clock                #    0.006 CPUs utilized
       3      context-switches          #   142.765 /sec
       0      cpu-migrations            #    0.000 /sec
      66      page-faults              #    3.141 K/sec
<not counted>  cpu_atom/cycles/                (0.00%)
  75,191,711    cpu_core/cycles/                #    3.578 GHz
<not counted>  cpu_atom/instructions/           (0.00%)
 162,384,235    cpu_core/instructions/           (0.00%)
<not counted>  cpu_atom/branches/              (0.00%)
  29,871,537    cpu_core/branches/              #    1.422 G/sec
<not counted>  cpu_atom/branch-misses/         (0.00%)
    114,826    cpu_core/branch-misses/
  TopdownL1 (cpu_core)                #   25.5 % tma_backend_bound
                                           #    3.6 % tma_bad_speculation
                                           #   32.9 % tma_frontend_bound
                                           #   38.0 % tma_retiring

  3.811052364 seconds time elapsed

  0.003914000 seconds user
  0.017879000 seconds sys
```

Multi-threaded (n=10) client connecting to **server using select**

CPU clocks: Moderate as the server efficiently manages multiple clients in one thread.

Cache misses: Relatively low compared to the concurrent model since only one thread is active.

Context switches: Low due to the absence of thread switching.

Latency: Comparable to the concurrent model since multiple clients are handled in parallel, but without the overhead of thread creation.

Performance counter stats for './multi_thread_client 25':

2.70 msec	task-clock	#	0.001 CPUs utilized	
56	context-switches	#	20.770 K/sec	
0	cpu-migrations	#	0.000 /sec	
114	page-faults	#	42.281 K/sec	
<not counted>	cpu_atom/cycles/			(0.00%)
8,948,927	cpu_core/cycles/	#	3.319 GHz	
<not counted>	cpu_atom/instructions/			(0.00%)
8,101,875	cpu_core/instructions/			
<not counted>	cpu_atom/branches/			(0.00%)
1,470,989	cpu_core/branches/	#	545.570 M/sec	
<not counted>	cpu_atom/branch-misses/			(0.00%)
26,965	cpu_core/branch-misses/			
TopdownL1 (cpu_core)		#	24.5 % tma_backend_bound	
		#	5.9 % tma_bad_speculation	
		#	45.8 % tma_frontend_bound	
		#	23.7 % tma_retiring	

1.980310343 seconds time elapsed

0.000000000 seconds user

0.002568000 seconds sys

Multi-threaded (n=25) client connecting to server using select

Performance counter stats for './select_server':

57.03 msec	task-clock	#	0.012 CPUs utilized	
10	context-switches	#	175.336 /sec	
0	cpu-migrations	#	0.000 /sec	
66	page-faults	#	1.157 K/sec	
<not counted>	cpu_atom/cycles/			(0.00%)
187,900,490	cpu_core/cycles/	#	3.295 GHz	
<not counted>	cpu_atom/instructions/			(0.00%)
414,647,214	cpu_core/instructions/			
<not counted>	cpu_atom/branches/			(0.00%)
76,344,414	cpu_core/branches/	#	1.339 G/sec	
<not counted>	cpu_atom/branch-misses/			(0.00%)
283,094	cpu_core/branch-misses/			
TopdownL1 (cpu_core)		#	24.7 % tma_backend_bound	
		#	3.2 % tma_bad_speculation	
		#	33.7 % tma_frontend_bound	
		#	38.4 % tma_retiring	

4.874951254 seconds time elapsed

0.017470000 seconds user

0.040060000 seconds sys

Multi-threaded (n=25) client connecting to server using select

```

Performance counter stats for './multi_thread_client 50':

      6.36 msec task-clock                #    0.000 CPUs utilized
        109    context-switches          #    17.138 K/sec
           0    cpu-migrations            #     0.000 /sec
        168    page-faults               #    26.414 K/sec
<not counted>    cpu_atom/cycles/                (0.00%)
19,072,763    cpu_core/cycles/                #     2.999 GHz
<not counted>    cpu_atom/instructions/          (0.00%)
14,967,477    cpu_core/instructions/          (0.00%)
<not counted>    cpu_atom/branches/              (0.00%)
2,718,206    cpu_core/branches/                #    427.369 M/sec
<not counted>    cpu_atom/branch-misses/          (0.00%)
40,609    cpu_core/branch-misses/
      TopdownL1 (cpu_core)                #    33.9 % tma_backend_bound
                                           #     4.2 % tma_bad_speculation
                                           #    40.6 % tma_frontend_bound
                                           #    21.2 % tma_retiring

39.138830256 seconds time elapsed

      0.000373000 seconds user
      0.005143000 seconds sys

```

Multi-threaded (n=50) client connecting to server using select

```

Performance counter stats for './select_server':

    125.29 msec task-clock                #    0.003 CPUs utilized
         16    context-switches          #    127.708 /sec
           0    cpu-migrations            #     0.000 /sec
         66    page-faults               #   526.794 /sec
<not counted>    cpu_atom/cycles/                (0.00%)
368,627,923    cpu_core/cycles/                #     2.942 GHz
<not counted>    cpu_atom/instructions/          (0.00%)
823,383,633    cpu_core/instructions/          (0.00%)
<not counted>    cpu_atom/branches/              (0.00%)
151,615,259    cpu_core/branches/                #     1.210 G/sec
<not counted>    cpu_atom/branch-misses/          (0.00%)
550,252    cpu_core/branch-misses/
      TopdownL1 (cpu_core)                #    23.6 % tma_backend_bound
                                           #     3.2 % tma_bad_speculation
                                           #    34.2 % tma_frontend_bound
                                           #    39.0 % tma_retiring

44.693547888 seconds time elapsed

      0.033495000 seconds user
      0.092365000 seconds sys

```

Multi-threaded (n=50) client connecting to **server using select**

CPU clocks: Increased but still lower than the concurrent model, as only one thread is used.

Cache misses: Slightly higher due to more connections being handled, but still lower than the multi-threaded approach.

Context switches: Very low as it is a single-threaded event loop.

Latency: Scales well with the number of clients but may show increased delay if too many clients are connected, as the server has to juggle multiple file descriptors in the event loop.

CPU utilization: Efficient use of CPU with lower overhead than the concurrent model, but may struggle under extremely high loads compared to a fully multi-threaded approach.

CONCLUSIONS:

Single-threaded: Best for small client loads with minimal overhead but scales poorly.

Concurrent: Efficient for high client loads but introduces significant overhead due to thread management and context switching.

Select-based: Offers a balance between single-threaded simplicity and concurrent parallelism, excelling for moderate loads with lower overhead but can struggle with very high client numbers.